

Pruning Neural Networks with Lottery Tickets in a MDP Approach

Andrey de Aguiar Salvi
Automated Planning 2019/02

From Previous Presentation...

We propose to modify the Lottery Tickets Hypothesis (a model compression method) with a Markov Decision Process with Q-Learning.

From Previous Presentation...

The Lottery Tickets Hypothesis.


Algorithm 1 Lottery Tickets Hypothesis

Source: adapted from (Frankle et al. 2019)

Require: weight matrix W , mask W' , pruning rate p

- 1: $W_0 \leftarrow W$
 - 2: $train(W, X, Y, n_epochs)$
 - 3: **while** $remaining_weights > total_weights \times p$ **do**
 - 4: $indexes \leftarrow find_smallest_values(|W|, p)$
 - 5: $W'[indexes] \leftarrow 0$
 - 6: $W \leftarrow W_0$
 - 7: $train(W, W', X, Y, n_epochs)$
 - 8: **end while**
-

Markov Decision Process



Source: Sutton, R. S., Barto, A. G., & Bach, F. (1998). Reinforcement learning: An introduction

Q-Learning

```
Q-learning
loop
   $a \leftarrow \text{Select}_a\{Q(s, a)\}$ 
  apply action  $a$ 
  observe resulting reward  $r(s, a)$  and next state  $s'$ 
   $Q(s, a) \leftarrow Q(s, a) + \alpha[r(s, a) + \max_{a'}\{Q(s', a')\} - Q(s, a)]$  (i)
   $s \leftarrow s'$ 
until termination condition
```

The Proposed Method

Algorithm 3 Our proposed method

Require: weight matrix W , mask W' , pruning rate p

- 1: $W_0 \leftarrow W$
 - 2: $train(W, X, Y, n_epochs)$
 - 3: **while** $remaining_weights > total_weights \times p$ **do**
 - 4: $s \leftarrow W'$
 - 5: $indexes \leftarrow \pi(s)$
 - 6: $W'[indexes] \leftarrow 0$
 - 7: $W \leftarrow W_0$
 - 8: $train(W, W', X, Y, n_epochs)$
 - 9: **end while**
-

Some Details

- Creating the Q-Table:

- First train LeNet 300-100 by 30 epochs
- The agent will perform a decreasing epsilon-search

- Total of steps: 1000
- Number of iterations: 15
- epsilon decreation [1.0; ...; 0.1]
- Reward: $r_a(s, s') \leftarrow \frac{-(1 - accuracy) * remaining_weights}{total_weights}$

- Q-Table update: or
$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r_a(s, s') + \max_{a'} \{Q(s', a')\} - Q(s, a) \right]$$
$$Q(s, a) \leftarrow (1 - \alpha) \cdot Q(s, a) + \alpha \left(r_a(s, s') + \gamma \cdot \max_a Q(s', a) \right)$$

Some Details

- Pruning the model:
 - Create another LeNet 300-100 from scratch
 - Iteratively:
 - Train 9 epochs
 - Prune (action which maximizes the Quality based on the state)

Results

Table 1: Train and validation accuracy's

Q-update	α	γ	Train_Acc	Valid_Acc
EQ_3	0.5	-	89.68%	89.31%
EQ_3	0.6	-	90.70%	90.40%
EQ_3	0.7	-	91.09%	90.86%
EQ_3	0.8	-	8.85%	0.09%
EQ_3	0.9	-	90.59%	89.82%
EQ_3	1.0	-	91.21%	90.86%
EQ_6	0.9	0.4	9.10%	0.09%
EQ_6	0.9	0.5	88.98%	88.69%
EQ_6	0.9	0.6	90.60%	89.94%
EQ_6	0.9	0.7	89.83%	89.22%
EQ_6	0.9	0.8	90.09%	89.39%
EQ_6	0.9	0.9	90.16%	89.79%
EQ_6	0.9	1.0	9.10%	0.09%
LTH	-	-	96.86%	95.92%

Results

Table 2: Test accuracy of the best models

Class	Model 1	Model 2	Model 3	Instances
0	97.44%	97.55%	98.57%	980
1	97.26%	97.18%	98.23%	1135
2	91.08%	88.56%	95.54%	1032
3	88.81%	87.62%	96.33%	1010
4	92.46%	94.80%	95.92%	982
5	83.85%	81.83%	94.05%	892
6	94.67%	94.05%	96.97%	958
7	91.34%	90.66%	96.39%	1028
8	89.73%	86.34%	94.76%	974
9	88%	86.91%	94.84%	1009
Overall	91.6%	90.7%	96.2%	10000

Thank you!

Andrey de Aguiar Salvi

andrey.salvi@edu.pucrs.br / andreysalvi@gmail.com